

# Linux Overview

**Advanced Embedded Linux  
Development  
with Dan Walkes**



University of Colorado **Boulder**

**Learning objectives:**

**Linux History**

**Linux Requirements**

**Linux contributors**

# Starting Out with Linux

- Embedded Linux started with the TiVo in 1999
- 2 billion devices running Linux as of 2017
- Why run Linux?
  - Moore's Law



# Points Driving Linux Adoption

- Functionality
  - Support built in for scheduler, network stack, USB, WiFi, Bluetooth, storage, etc
- Ported to wide range of architectures.
- Open source, modifiable
- Active community - answer questions.
- No vendor lock-in

# Linux Considerations

- Needs a 32 bit processor and ~16+MB of RAM, ~8MB of flash
  - Needs Memory Management Unit for all practical purposes
- Needs skill set of engineers (IE you!)
- May not be appropriate for some real-time applications.

# The Players

- Open Source Community
  - Alliance of developers, not for profit, academic, commercial
  - Group for each set of applications
- CPU Architects/SOC Vendors/Board Vendors
  - Create reference hardware/ board support packages (BSP)s

# Linux Command Line

**Advanced Embedded Linux  
Development  
with Dan Walkes**

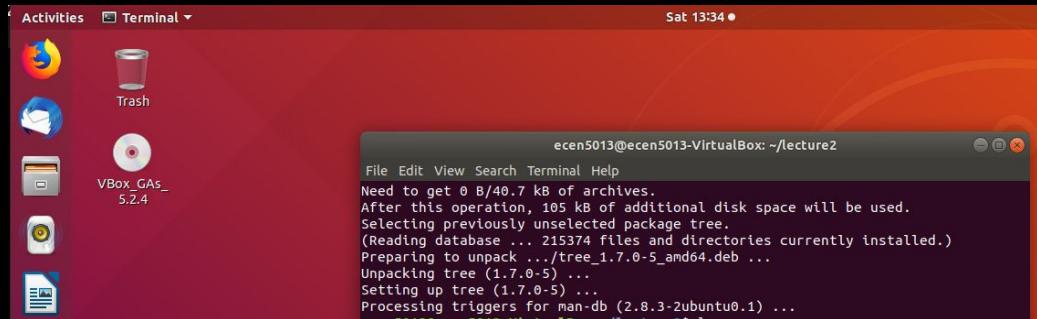


University of Colorado **Boulder**

**Learning objectives:**  
**Overview of Linux Command Line**  
**Common Linux Utilities**

# Why use the command line?

- Ubuntu has a full graphical UI... can I use that?



- Absolutely!
- However, our assignments are going to focus on automation and scripting - which require the terminal

# Why use the command line?

- The command line is the only thing <every> linux system has.
  - This includes your embedded device.
- It's often the only thing you have when troubleshooting problems.
- It's often the easiest, fastest, and most reliable way to remotely access a system.

# Common Linux Utilities - I/O

- echo
  - echo a command to the terminal

```
ecen5013@ecen5013-VirtualBox:~/lecture2$ echo "Hello World!"  
Hello World!
```

- cat
  - concatenate (print out) a file to the terminal

```
ecen5013@ecen5013-VirtualBox:~/lecture2$ cat myfile.txt  
ECEN_5013_IS_AWESOME!
```

# Common Linux Utilities - Directory Navigation

- `pwd`
  - Print Working Directory

```
ecen5013@ecen5013-VirtualBox:~$ pwd  
/home/ecen5013
```

- `ls`
  - List contents of a directory

```
ecen5013@ecen5013-VirtualBox:~$ ls  
aesd-assignments assignments-complete-private Desktop Documents Downloads ecen5013-buildroot ecen5013-source ecen5013-yocto examples.desktop Music Pictures Public Templates Videos
```

# Common Linux Utilities - Directory Navigation

- cd
  - Change directory

```
ecen5013@ecen5013-VirtualBox:~$ cd aesd-assignments/
ecen5013@ecen5013-VirtualBox:~/aesd-assignments$ ls
LICENSE README.md tester.sh
```

- mkdir
  - Make a new directory

# Common Linux Utilities - Directory Navigation

- mv
  - Move files or directories to a new location
- cp
  - Copy files or directories to a new location
- rm
  - Remove a directory

# Common Linux Utilities

- touch
  - Update a file's timestamp without editing if the file exists.
  - Create an empty file if the file does not exist.

# Common Linux Utilities - man

- man < program or command>
  - “man pages” (short for manual pages)
  - Dates back to the Unix Programmer's Manual, 1971

```
CP(1)                               User Commands                               CP(1)

NAME
    cp - copy files and directories

SYNOPSIS
    cp [OPTION]... [-T] SOURCE DEST
    cp [OPTION]... SOURCE... DIRECTORY
    cp [OPTION]... -t DIRECTORY SOURCE...

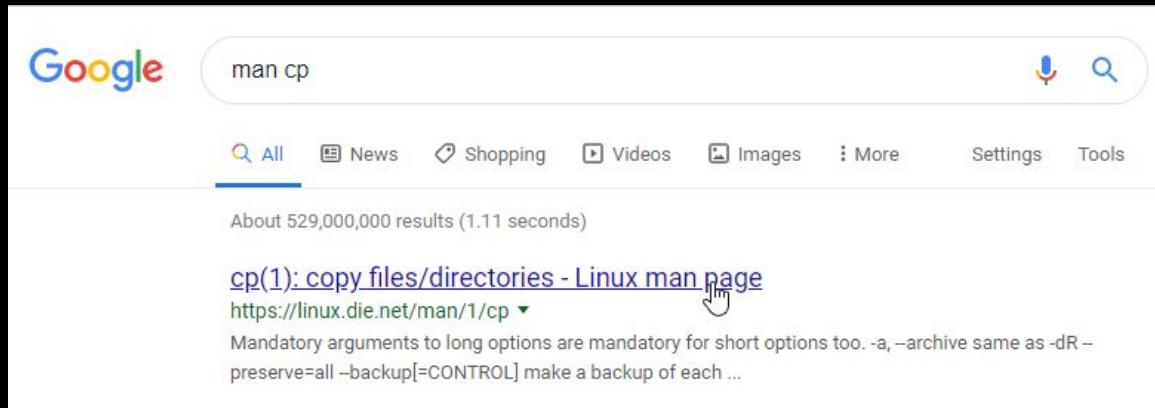
DESCRIPTION
    Copy SOURCE to DEST, or multiple SOURCE(s) to DIRECTORY.

    Mandatory arguments to long options are mandatory for short options
    too.

    -a, --archive
        same as -dR --preserve=all
```

# Common Linux Utilities - man

- Google knows about man pages too!



# Common Linux Utilities - man

- Example of common copy command issue

```
dan@DESKTOP-BQMV69:~/CU/aesd-lectures/lecture2$ mkdir a_directory
dan@DESKTOP-BQMV69:~/CU/aesd-lectures/lecture2$ mkdir a_directory/a_subdir
dan@DESKTOP-BQMV69:~/CU/aesd-lectures/lecture2$ mkdir a_copy_target
dan@DESKTOP-BQMV69:~/CU/aesd-lectures/lecture2$ cp a_directory/ a_copy_target/
cp: -r not specified; omitting directory 'a_directory'
```

## cp(1) - Linux man page

### Name

cp - copy files and directories

### **-R, -r, --recursive**

copy directories recursively

```
dan@DESKTOP-BQMV69:~/CU/aesd-lectures/lecture2$ cp -r a_directory/ a_copy_target/
```

# cp -r

- Fix to use cp -r, validate with tree

```
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ cp -r a_directory/ a_copy_target/
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ tree a_copy_target/
a_copy_target/
└── a_directory
    └── a_subdir

2 directories, 0 files
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ tree a_directory/
a_directory/
└── a_subdir

1 directory, 0 files
```

# Advanced Command Line

**Advanced Embedded Linux  
Development  
with Dan Walkes**



University of Colorado **Boulder**

**Learning objectives:**

**Editing from command line**

**Using the shell**

**Searching from command line**

**Wildcards, pipes, redirection**

**Permissions**

**Remote access to the command line**



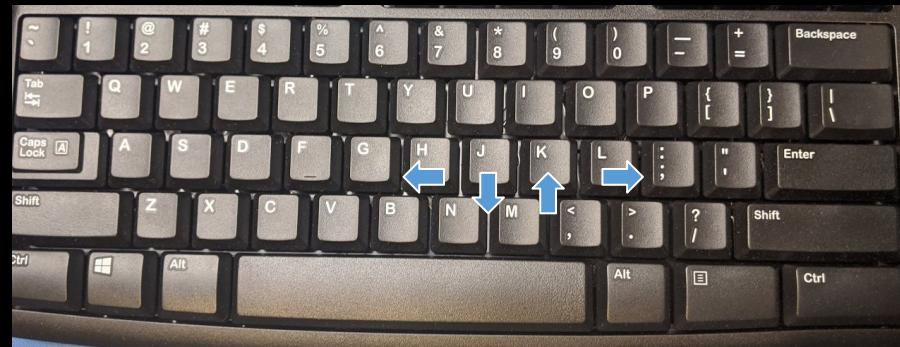
# Editing - vi

- Linux has lots of text editors to choose from.
- “vi” is one classic utility
  - Short reference for the “visual” mode of an early text editor utility on UNIX
  - Written in 1976
  - Often replaced by “vim” “VI improved”
- I suggest getting at least some familiarity with “vi”
  - Why? - It’s everywhere. In my experience it’s the only editor you can always count on being available.

# Editing - vi

- Basic commands:
  - vi <name> (or vim <name>)
- Entering “normal mode” for commands - use Esc key
  - write - Esc -> :w Esc -> :q - quit
  - Insert: i, Delete d
  - Navigation keys

Move	Key
Left	h
Down	j
Up	k
Right	l



# Shell Interpreter Options

- A shell interpreter handles commands from the terminal (or from a script file).
- Several have been developed, many have similar command support.
- Bourne shell (sh) was one of the originals, released in 1979
- Bash shell (Bourne Again SHell) was created as a replacement for Bourne with features from ksh and csh.

# Searching for Files

- Locate files/directories with find
  - Find files/directories below the current directory  
“.” with name a\_directory

```
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ find . -name a_directory
./a_copy_target/a_directory
./a_directory
```

# Searching for Content

- Searching for content with grep

```
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ cat myfile.txt
AESD_IS_AWESOME!
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ cat myfile2.txt
CU_IS_AWESOME!
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ grep -r "IS_AWESOME" *
myfile.txt:AESD_IS_AWESOME!
myfile2.txt:CU_IS_AWESOME!
```

# Wildcards

```
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ cat myfile.txt  
AESD_IS_AWESOME!  
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ cat myfile2.txt  
CU_IS_AWESOME!  
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ grep -r "IS_AWESOME" *  
myfile.txt:AESD_IS_AWESOME!  
myfile2.txt:CU_IS_AWESOME!
```

- \* in means files with any character in the name
- myfile\*txt would have meant any file starting with myfile and ending with txt
- expanded by the shell to
  - grep -r “IS\_AWESOME” myfile.txt myfile1.txt myfile2.txt myfile3.txt myfile4.txt

# Pipes

- The Pipe Character ‘|’ sends the output of one command to the input of another command
  - Chain commands together based on standard input/output streams.

```
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ cat myfile.txt | grep "IS_AWESOME"  
AESD_IS_AWESOME!
```

# Redirection

- Send the output of a command to a new file using >

```
dan@DESKTOP-BQMV69:~/CU/aesd-lectures/lecture2$ grep -r "IS_AWESOME" * > searchresult.txt
dan@DESKTOP-BQMV69:~/CU/aesd-lectures/lecture2$ cat searchresult.txt
textfile.txt:AESD_IS_AWESOME!
textfield2.txt:CU_IS_AWESOME!
```

- Append the output of a command into an existing file using >>

```
dan@DESKTOP-BQMV69:~/CU/aesd-lectures/lecture2$ echo "hello" >> searchresult.txt
dan@DESKTOP-BQMV69:~/CU/aesd-lectures/lecture2$ cat searchresult.txt
textfile.txt:AESD_IS_AWESOME!
textfield2.txt:CU_IS_AWESOME!
hello
```

# File Permissions

- Control how a file may be used and by whom
- 3 levels of permission - User, Group and World (or everyone)

Use ls -l to show permission info

User (owner)			Group			Others (everyone)		
Read (r)	Write (w)	Execute (x)	Read (r)	Write (w)	Execute (x)	Read (r)	Write (w)	Execute (x)

```
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ ls -l
total 20
drwxr-xr-x 3 dan dan 4096 Jun 21 15:28 a_copy_target
drwxr-xr-x 3 dan dan 4096 Jun 21 15:28 a_directory
-rw-r--r-- 1 dan dan   65 Jun 21 15:35 searchresult.txt
-rw-r--r-- 1 dan dan   17 Jun 21 15:31 textfile.txt
-rw-r--r-- 1 dan dan   15 Jun 21 15:31 textfile2.txt
```

# File Permissions

- Execute permissions typically not granted by default

```
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ cat helloworld.sh
#!/bin/bash
echo "hello world!"
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ ls -l helloworld.sh
-rw-r--r-- 1 dan dan 32 Jun 21 15:38 helloworld.sh
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ ./helloworld.sh
-bash: ./helloworld.sh: Permission denied
```

- Use chmod to change permissions

```
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ chmod u+x helloworld.sh
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ ls -l helloworld.sh
-rwxr--r-- 1 dan dan 32 Jun 21 15:38 helloworld.sh
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ ./helloworld.sh
hello world!
```

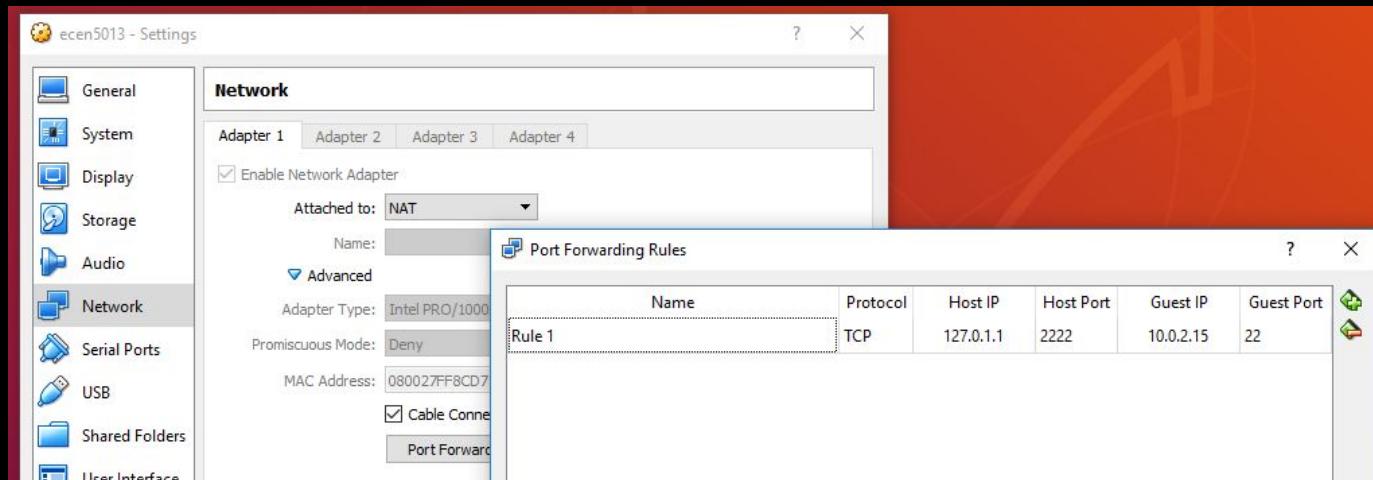
# File Permissions

- Permissions can be specified by octal digits in numeric mode.
  - 4 = read
  - 2 = write
  - 1 = execute

```
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ ls -l helloworld.sh
-rwxr--r-- 1 dan dan 32 Jun 21 15:38 helloworld.sh
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ chmod 766 helloworld.sh
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ ls -l helloworld.sh
-rwxrw-rw- 1 dan dan 32 Jun 21 15:38 helloworld.sh
```

# Remote Access

- You can use Secure Shell (ssh) for remote access to machines (including your VM)



```
danwa@DESKTOP-BQMVP69 MSYS ~  
$ ssh -p 2222 dan@127.0.1.1|
```

# Remote Access

- Use a unix environment WSL or msys2 on Windows
  - <https://docs.microsoft.com/en-us/windows/wsl/install-win10>
  - <https://www.msys2.org/>
- Use scp for file transfer into/out of your VM

Transfer from windows to your virtual machine home directory

```
danwa@DESKTOP-BQMVP69 MSYS ~
$ scp -P 2222 /c/Users/danwa/Downloads/somefile.txt dan@127.0.1.1:~|
```

SSH into your virtual machine

```
danwa@DESKTOP-BQMVP69 MSYS ~
$ ssh -p 2222 dan@127.0.1.1|
```

Verify your virtual machine contains somefile.txt

```
dan@DESKTOP-BQMVP69:~$ ls -la | grep somefile
-rw-r--r-- 1 dan dan 0 Jun 21 15:45 somefile.txt
```

Transfer from virtual machine into your Windows system

```
danwa@DESKTOP-BQMVP69 MSYS ~
$ scp -P 2222 dan@127.0.1.1:~/somefile.txt /c/Users/danwa/Downloads/from_vm.txt
```

# Scripting

## Advanced Embedded Linux Development with Dan Walkes



University of Colorado **Boulder**

**Learning objectives:**

**Scripts Overview**

**Script Variables**

**Script Arguments**

**Script Conditionals**



# Scripts

- Scripting allows us to:
  - Automate frequently used commands
  - Customize commands for different scenarios using arguments and variables
- Think of scripts as documentation!
  - If you find yourself listing commands used to perform a task, consider scripting instead.
- Often are in files ending with .sh (short for shell)

# Simple Script Example

- `#!/bin/sh` is a comment which tells the shell which shell interpreter to use
  - Also known as “shebang”
- Also common to see `#!/bin/bash`

```
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ cat helloworld.sh
#!/bin/bash
greeting=hello
echo "${greeting} world!"
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ ./helloworld.sh
hello world!
```

# Script Variables

- Script Variables include:
  - Variables we define ourselves

```
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ cat helloworld.sh
#!/bin/bash
greeting=hello
echo "${greeting} world!"
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ ./helloworld.sh
hello world!
```

- Variables passed into the script as arguments

```
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ cat helloworldargs.sh
#!/bin/bash
greeting=$1
echo "${greeting} world!"
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ ./helloworldargs.sh hola
hola world!
```



# Script Argument

- Script special variables
  - \$0 - the invoked name of the bash script
  - \$1 - \$9 - first 9 arguments to a script
  - \$# - Number of arguments passed to the script
  - \$? - The exit status of the most recent process

# Bash Conditional if/else

- Conditional syntax for bash
  - [ ] is shorthand for the “test” command

if [ conditional ]  
then

action  
else

action

```
if [ -d "$WRITEDIR" ]
then
    echo "$WRITEDIR created"
else
    exit 1
fi
```

Digitized by srujanika@gmail.com

-d FILE

FILE exists and is a directory

# Script Failure Exit Codes

- Scripts can communicate success/failure through exit status and the exit command

```
if [ $? -eq 0 ]; then
    echo "success"
    exit 0
else
    echo "failed: expected ${MATCHSTR} in ${OUTPUTSTRING} but instead found"
    exit 1
fi
```

```
dan@DESKTOP-BQMVP69:~/CU/assignments-complete-private$ ./finder.sh
ERROR: Invalid Number of Arguments.
Total number of arguments should be 2.
The order of the arguments should be:
    1)File Directory Path.
    2)String to be searched in the specified directory path.
```

# Course Introduction

**Advanced Embedded Linux  
Development  
with Dan Walkes**



University of Colorado **Boulder**

**Learning objectives:**

**Introduction to the course**

**Why learn Linux?**

**What this course is and isn't**

# Who Is This Guy?

- Dan Walkes
- Work full time in Industry and teach part time
  - VP of Embedded Engineering at Sighthound
- Began teaching Spring 2019 semester
- Former CU Boulder ESE student in the early 2000s
  - Former TA for ECEN 5623 Real Time Embedded

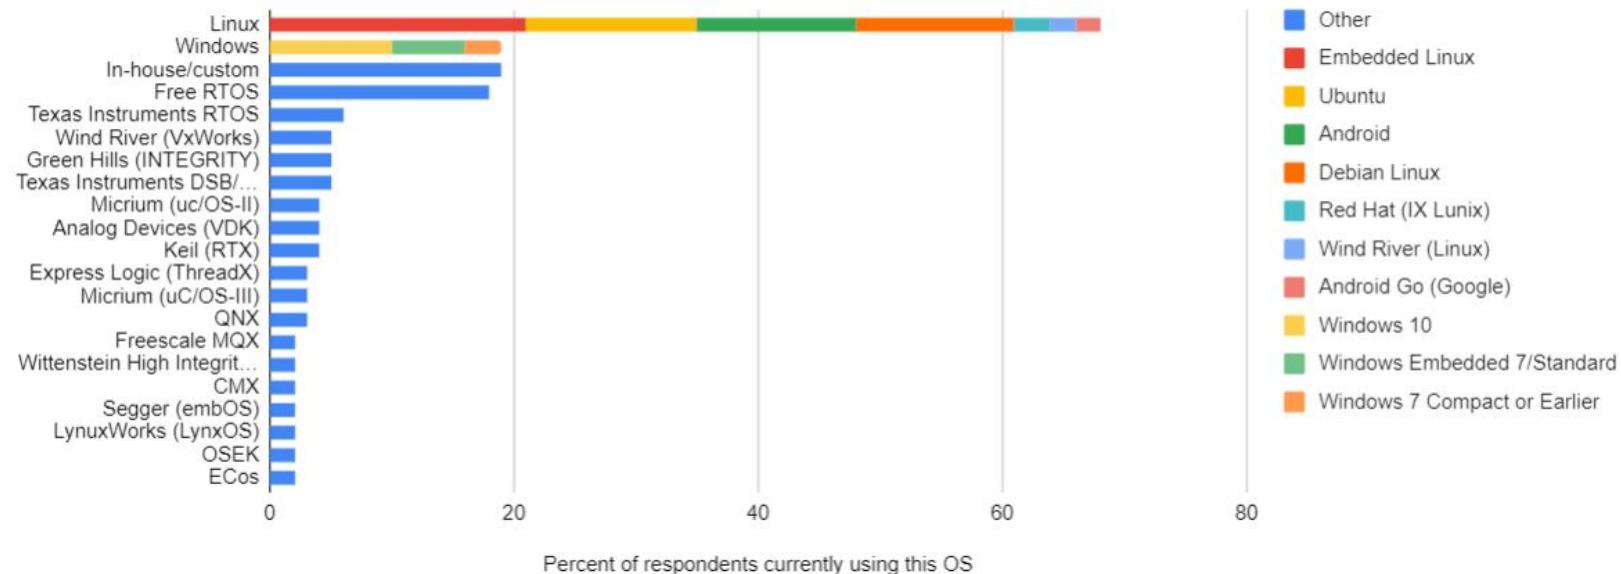


# Who Is This Guy?

- 20th year of Industry experience in 2019
  - Most of the past 10 has been Linux embedded systems related.
- Not a Dr. or Professor. Please call me Dan.
  - Think of me as a team lead on your development team.

# Why Learn Linux?

2019 Embedded Markets Study Operating Systems



- Based on 958 responses Jan-March 2019

# Course Objectives

- Understand how to configure and deploy a Linux based Embedded System.
- Gain experience with common Linux system build tools and components:
  - BASH and shell scripting
  - Buildroot and Yocto
  - QEMU

# Course Objectives

- Gain experience with Linux Driver development
- Improve your resume/interview skills in relation to Embedded Linux!
- Cover topics I wish I would have had as a part of my graduate studies.
  - In many cases they didn't exist yet.

# What This Class is NOT

- Not a data structures or algorithms course.
  - We will cover these only from the standpoint of how to use existing libraries/software.
- Not a low resource microcontroller/FreeRTOS course.
  - Not covering event loops or protocols like Bluetooth
  - Not specifically covering sensors
  - Not covering low resource embedded libraries.

# Course Format

**Advanced Embedded Linux  
Development  
with Dan Walkes**



University of Colorado **Boulder**

**Learning objectives:**

**Prerequisites and expectations**

**Assignments and build environments**

**Books**



# Prerequisites and Expectations

- I expect some working knowledge of:
  - C Programming
  - Makefiles
- Ideally you'd have some knowledge of:
  - Linux command line
  - Shell scripting

# Assignments and Build Environment

- A Linux build environment will be required for all assignments.
  - This can be a virtual machine or actual Linux hardware.
  - See Assignment 1 instructions for version in use.
- Option 1: Use VirtualBox on Windows or MacOS
  - Install a Ubuntu VM with at least 4GB of RAM and 100GB of disk space
  - Requires ~8GB of RAM and 100GB of host disk space
  - Some students use a USB 3.0 HDD/SSD for VM images

# Assignments and Build Environment

- Option 2: Use a dedicated Linux laptop
  - Find one with at least 8GB of RAM and 100GB of disk space.
  - Doesn't need to be expensive: think pawn shops and ebay.
  - It's a good idea to google <hardware model> Ubuntu to see if you can find people stating it's supported.
- Option 3: Use a cloud hosted runner

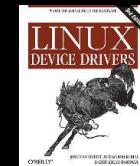
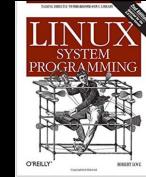


# Assignments / Project and git

- We will make extensive use of git in this course.
  - You will use git to turn in all assignments through Github Classroom.
- Assignments may start with git merge.
- You may see merge issues/errors. When you do:
  - Use an internet search! Lots of help is available online
  - Use appropriate help channels

# Books

- **Linux System Programming, 2nd Edition (LSP)**
  - How your software programs will interact with the kernel.
  - Processes, Threading, POSIX
- **Mastering Embedded Linux Programming, 2nd Edition (MELP)**
  - Practical aspects of building Embedded Systems
- **Linux Device Drivers 3rd Edition (LDD)**
  - Theory and practice related to Linux Device Driver development.



# Environment Setup

**Advanced Embedded Software  
Development  
with Dan Walkes**



University of Colorado **Boulder**

## **Learning objectives:**

**Introduction to Linux Workshop**

**Installing a Linux Build VM**

**VM Snapshots**

**Distributions and Packages**

**Ubuntu and root access**

# Lecture Backup Material

- **Introduction to Linux Workshop**

<https://nsdl.oercommons.org/courses/an-introduction-to-linux-2/view>

- Shells, Text Editors (vim section), Remote Connections, Filesystem, File Permissions, Optional Topics

- **Bash Scripting Tutorials:**

<https://linuxconfig.org/bash-scripting-tutorial-for-beginners>

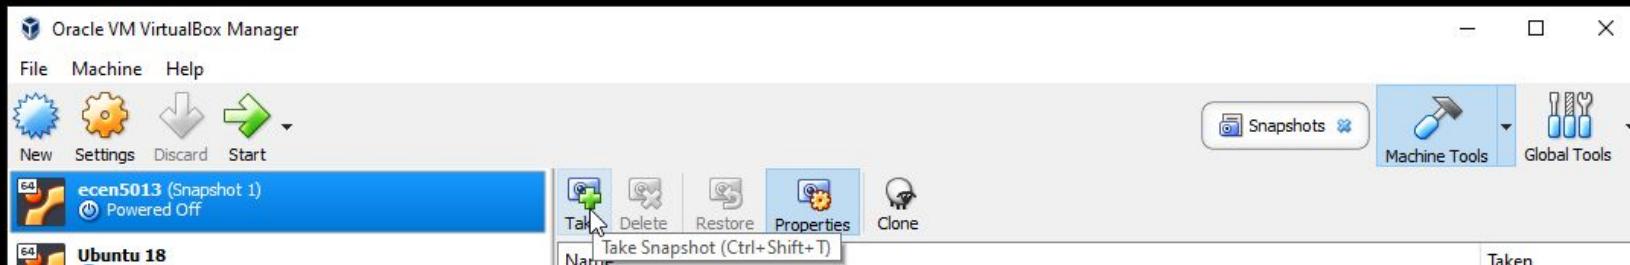
<https://ryanstutorials.net/bash-scripting-tutorial/bash-variables.php>

# Installing a Linux Build VM

- If you are running Windows or MacOS, we will use VirtualBox for our Linux based build environment.
  - See <https://www.virtualbox.org/>
- We will use VirtualBox + Ubuntu.
  - See the [Environment Setup](#) document for details.

# Virtual Machine Snapshots

- Use these in case your VM crashes, to recover your last known working configuration



# Using a Physical Linux Machine

- If your laptop is not powerful enough to run VirtualBox you can also run on dedicated Linux hardware.
  - Does not need to be expensive - lots on ebay or in pawn shops for < \$200
- You may also consider dual booting

# Distributions and Packages

- Linux is available in many distributions.
  - Distributions are collections of preconfigured software maintained by the community or companies.
- Software version and configuration is managed in Packages.
- Ubuntu is a popular Linux distribution.
- Each distribution has its own mechanism of adding software packages through the command line.

# Packages vs Binaries

- Why doesn't Linux just use binary installers like Windows does? Why use package management?
  - Packages were essentially the precursor to "app stores"
  - Package managers provide a centralized location for trusted software sources.
  - Also solves fragmentation problems
    - Software packages may need slight customizations for different distributions.

# What is Ubuntu?

- Pronounced “uu-boonto”
  - African word meaning “humanity to others”
- First release in 2004
- Based on another distribution known as “Debian”
  - Compared to Debian it’s (arguably) less stable but more up to date.
  - Ubuntu is more focused on usability than licensing.

# Ubuntu Package Manager

- Ubuntu's package management was traditionally performed by the “apt” utility
- Recent Ubuntu releases also support the “snap” utility
- Ubuntu will often suggest packages

```
ecen5013@ecen5013-VirtualBox:~/lecture2$ tree a_copy_target/
.
Command 'tree' not found, but can be installed with:
  sudo snap install tree  # version 1.8.0+pkg-3fd6, or
  sudo apt  install tree

See 'snap info tree' for additional versions.
```

# Linux Root user

- A default user account has access limitations.
  - Prevents you from accidentally deleting/modifying things you didn't intend.
    - example:
      - `rm -rf ${undefined_variable}`
- The “root” user can access/modify anything on the system
- Different permission levels allow the system administrator to configure appropriate permissions per user

# Ubuntu Root Access Philosophy

- By default your account is able to run commands as root using the “sudo” (short for super user do) command.
- sudo <command name>
  - Asks you for your password
  - Runs the command you specified as root

# System Programming Overview

**Advanced Embedded Linux  
Development  
with Dan Walkes**



University of Colorado **Boulder**

**Learning objectives:**  
**Introduce System Programming**  
**Understand APIs and ABIs**  
**Introduce POSIX**



# System Programming Overview

- What is System Software?
  - Interfacing with the kernel and C library.
  - GUI, compiler, debugger, web server, database.
- Differences relative to “Application Software”:
  - Doesn’t use higher level libraries.
  - Less OS/hardware details abstracted.



# Cornerstones of System Programming

- System Calls (syscalls)
- C Library (libc)
- C Compiler/linker (toolchain)



# Cornerstones of System Programming

- System Calls (syscalls)
  - Kernel function invocations from user space (read() write()) Roughly 300 total.
  - Shared subset of ~90% implemented by all architectures on Linux.

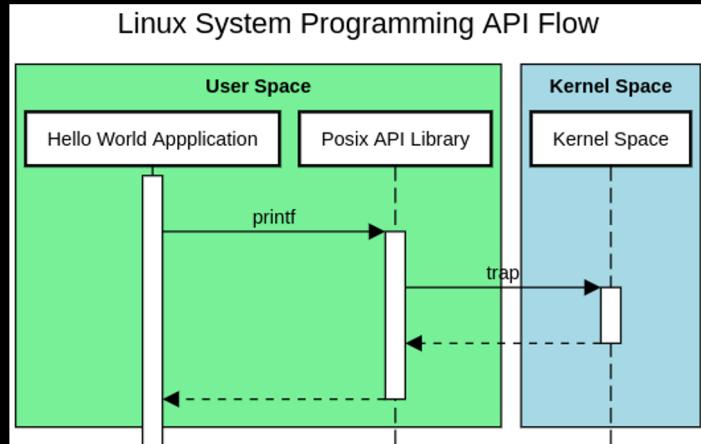


# Cornerstones of System Programming

- C Library (libc)
  - GNU libc (glibc) (gee-lib-see) Wrappers for system calls, threading support and applications
- C Compiler/linker (toolchain)
  - GCC (GNU C compiler)

# Invoking System Calls

- Not possible to link your user space application with the Linux kernel.
  - Why not?
    - Not allowed for security/reliability
- How do you invoke system calls?
  - Use a “trap” - typically a software interrupt





# API

- Application Programming Interface.
- API is a definition.
- Software which provides an API is the implementation.
- Ensures **source** compatibility.
  - Source can be compiled for different platforms, works in a specific way.

# API

- Source code remains portable across different hardware platforms/revisions (ideally)
- Example: C library functions like `printf()`, `strcpy()`, etc.
- Exceptions: `syscall` ~10% architecture differences mentioned in previous slide

# ABI

- Application Binary Interface
- Calling conventions, byte ordering, register use
- Ensures **binary** compatibility
- Defined/implemented by Kernel and toolchain
- Defines application interaction with itself, libraries, the kernel.

# ABI

- Binary generated after compilation/link (by toolchain).
- Byte code specific hardware types, software/compiler/library revisions.

Requires match:

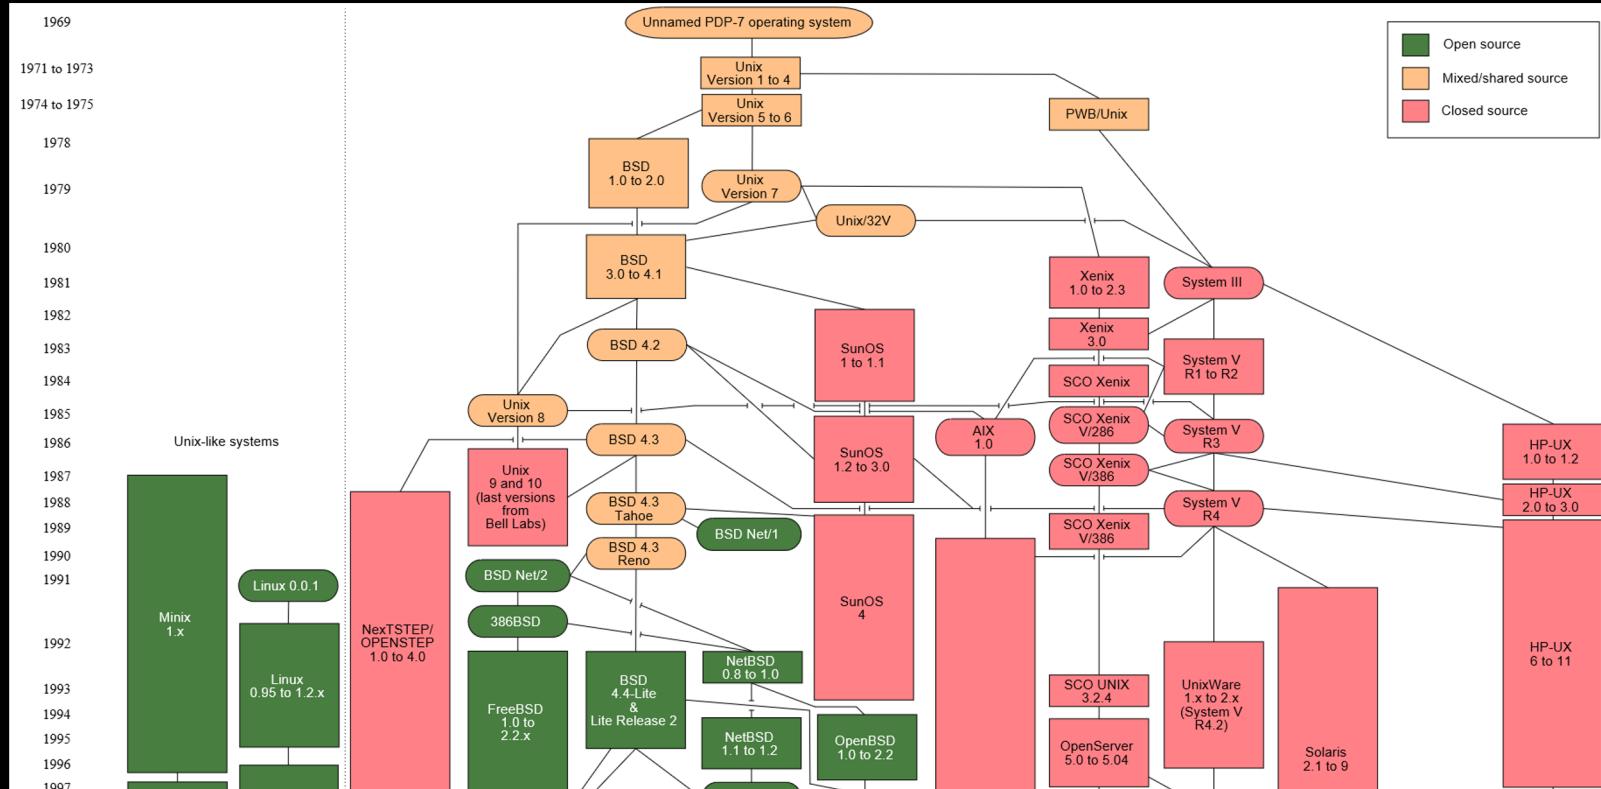
- hardware
- software libraries (especially glibc)
- compiler



# POSIX

- C APIs for Unix-like operating systems
- Stands for Portable Operating System Interface
  - Pronounced “pahz-icks”
- Started by IEEE in the late 1980s, as a way to coalesce the “Unix Wars”

# UNIX Wars



By Eraserhead1, Infinity0, Sav\_vas - Levenez Unix History Diagram, Information on the history of IBM's AIX on ibm.com, CC BY-SA 3.0,

# Enter Linux

Linus Benedict Torvalds 8/25/91

★ Hello everybody out there using minix -

I'm doing a (free) operating system (just a hobby, won't be big and professional like gnu) for 386(486) AT clones. This has been brewing since april, and is starting to get ready. I'd like any feedback on things people like/dislike in minix, as my OS resembles it somewhat (same physical layout of the file-system (due to practical reasons) among other things).

I've currently ported bash(1.08) and gcc(1.40), and things seem to work. This implies that I'll get something practical within a few months, and I'd like to know what features most people would want. Any suggestions are welcome, but I won't promise I'll implement them :-)

Linus ([torg...@kruuna.helsinki.fi](mailto:torg...@kruuna.helsinki.fi))

PS. Yes - it's free of any minix code, and it has a multi-threaded fs. It is NOT portable (uses 386 task switching etc), and it probably never will support anything other than AT-harddisks, as that's all I have :-.

# POSIX

- Open Software Foundation created the Single Unix Specification (SUS)
  - Specification was free
  - Eventually incorporated the POSIX standard.

# POSIX

- POSIX defines C code API for many concepts we'll cover this semester
  - File and directory operations
  - Clocks and timers
  - Semaphores
  - Shared memory
  - Threads

# POSIX

- Official list of headers:

<http://pubs.opengroup.org/onlinepubs/9699919799/idx/head.html>

- Some familiar ones:
  - <stdio.h>
  - <stdint.h>

# Linux Filesystems

**Advanced Embedded Linux  
Development  
with Dan Walkes**



University of Colorado **Boulder**

**Learning objectives:**

**Understand use of files in Linux**

**Introduction to Inodes**

**Introduction to Links**

**Linux filesystems**

# “Everything Is A File”

- Much of the interaction with the kernel occurs via reading/writing files.
- Devices are accessed similar to files
  - Example /dev/ttyUSB0 for USB serial port.
- Common open, manipulate, close paradigm shared with files and devices.

# Files in Linux

- Represented by paths.
  - Absolute “/path/to/file” - from root of filesystem.
  - Relative “path/to/file” - from working directory.
    - No leading “/” = relative path.

# Linux Regular File Properties

- Bytes of data in a byte stream.
  - No structure enforced at the system level.
- File position/offset tracks location.
- Can be truncated to size smaller or larger than original.

# Linux Regular File Properties

- Can be opened more than once, by different or same processes.
- Referenced and accessed by inode in the filesystem.

# Linux Inode (Index Node)

- Metadata used to track files on disk.
- Includes timestamp, owner, size, mode (access permission), location.
- Fixed small (128 byte) size.
  - Filename is not included.

Inode					
Accessed Time	Size	UID	GID	Permission	Location On Disk

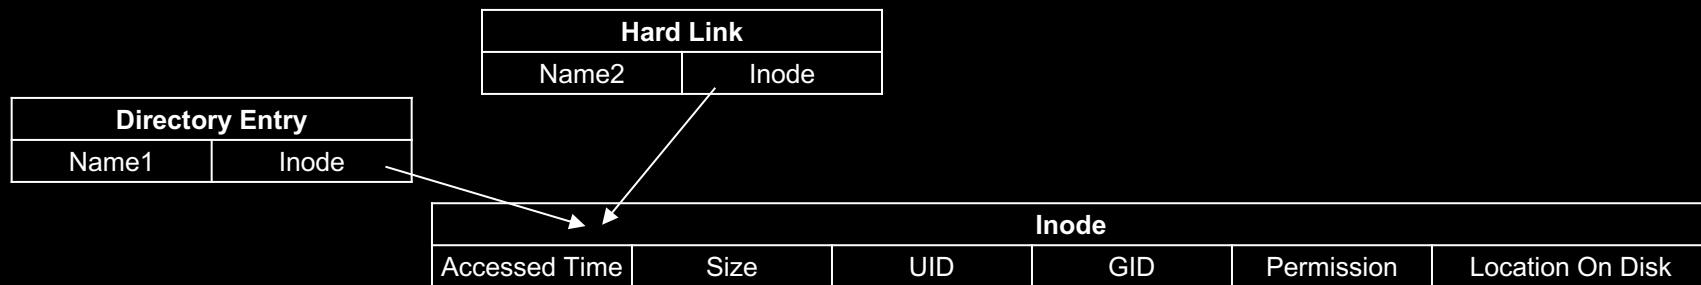
# Linux Directories

- Map human readable names to Inode numbers.
- Actually just files (with their own inodes) containing mapping of names to Inodes.
- Start at the root directory (“/”)



# Linux Directories

- Why not include filenames in Inodes?
  - Allows different file names to share the same content without duplicating Inode content.



# Linux Links

- Flexibility of Inode design means multiple names can resolve to the same Inode.
- Links redirect two file/directory paths to the same Inode

```
user@myhost:~/myfolder$ ls -l
total 4
-rw-r--r-- 1 user user 29 May  8 17:09 myfile.txt
lrwxrwxrwx 1 user user 10 May  9 05:55 myLink -> myfile.txt
```

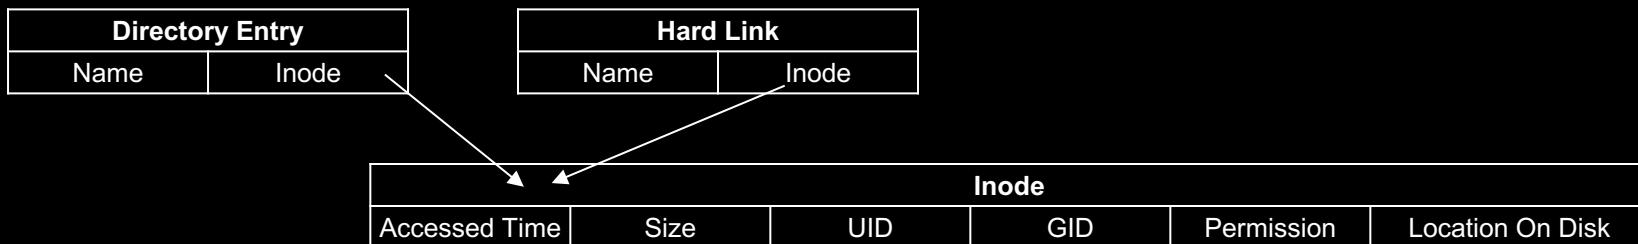


# Linux Links

- Two types of Links, Hard and Symbolic (symlinks)
  - Hard Links map directly to inodes, only allowed on the same filesystem.
  - Soft links map to filenames, work across filesystems, can be broken.

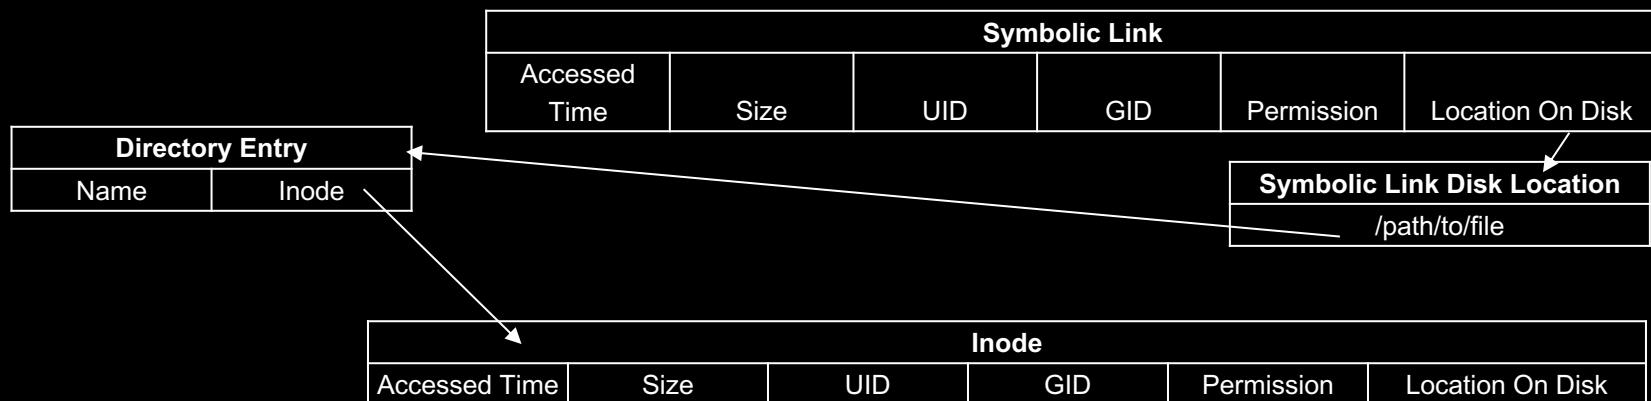
# Hard Links

- Can't span filesystems (Inode references a specific filesystem.)
- Conceptually the same as a Directory Entry.
- File deletes aren't allowed until all references are deleted (preventing broken Hard Links).



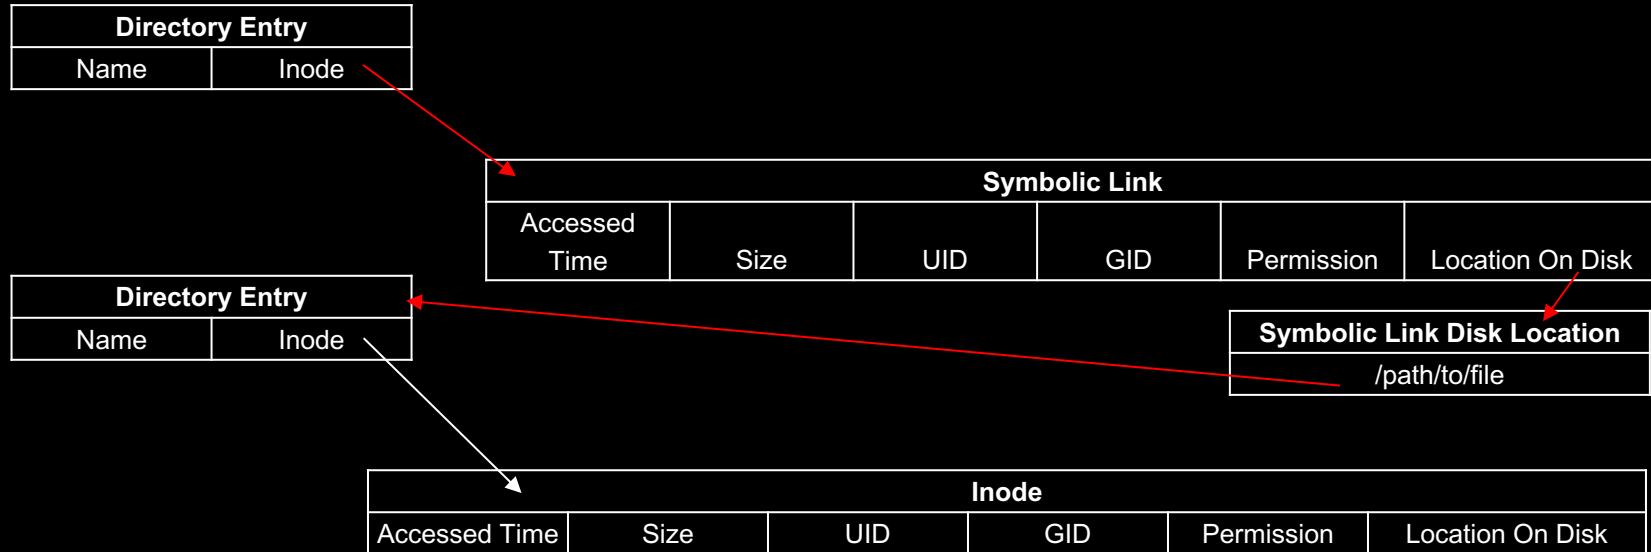
# Symbolic Links (symlinks)

- Regular file with complete path in content
- Can point anywhere, including other filesystems, can break



# Symbolic Links

- Requires extra steps to resolve





# Linux Special Files

- Map hardware devices to the “everything is a file” paradigm.
- Kernel objects represented as files.
  - Character Device
  - Block Device
  - Named Pipes
  - Sockets

# Linux Special Files

- Character Device
  - Linear queue of bytes
  - Example: Keyboard
- Block Device
  - Array of bytes, addressable in a sector
  - IE Hard Disk
- Named Pipes/Sockets
  - Interprocess Communication

# Linux Filesystem

- Collection of files in a hierarchy
- Specific types supported, tied to storage types
  - NFS - network file storage
  - ext4 - block device storage
  - fat - Microsoft defined storage format for disks

# Linux Filesystem

- Mounted/Unmounted to add/remove from the root filesystem.
- Smallest unit addressable is a block
  - block is a power of two multiple of sector size
  - Typical/historical sector size is 512 bytes

# Linux Blocks/Sectors

- Consider a 1TB Filesystem.
  - 1,000,000,000,000 bytes.
  - 0x3B98CCA00 bytes.
  - 34 bits required to address individually.

# Linux Blocks/Sectors

- 34 bits to address each byte.
- With 512 byte blocks:
  - Only 1,953,125,000 (0x746a5288) blocks.
  - 30 bits required to access each block.
- With 4K blocks:
  - Only 244,140,625 (0xE8D4A51) blocks.
  - 28 bits required to access each block.

# Processes and Threads

**Advanced Embedded Linux  
Development  
with Dan Walkes**



University of Colorado **Boulder**

**Learning objectives:**

**Understand Linux Processes**

**Understand Linux Threads**

**Introduce Interprocess  
Communication (IPC)**

# Linux Processes

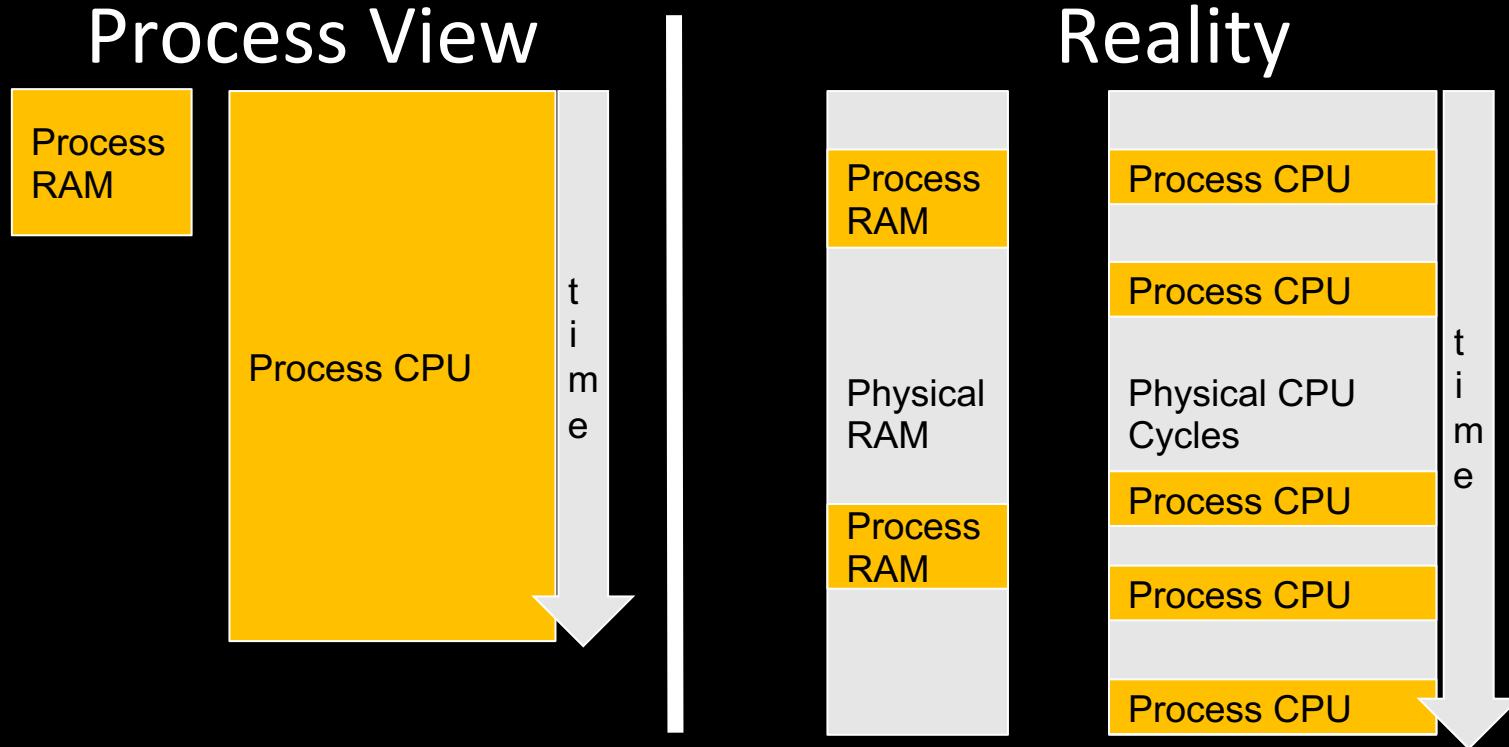
- Executable object code, running on hardware.
  - ELF (Executable and Linkable Format).
- Necessary resources to run are allocated and managed by the kernel through system calls.
  - Timers
  - Files
  - Hardware access



# Linux Processes

- Referenced by Process ID (pid).
- Run on virtualized processor and virtualized view of memory.
  - Appears to software as its own dedicated RAM/CPU.

# Virtualized Processor/Memory



# Linux Threads

- Unit of activity within a process
- A process may be single-threaded or multithreaded
- Each thread has
  - Stack - stores local variables
  - Processor state/current location
- **Memory address space is shared between threads**

# Linux Threads/Processes and Memory



- Each Process has its own virtual memory.
- Each Thread shares process virtual memory.
- Sharing memory access between threads?
  - Access directly (use synchronization).
- Sharing memory access between processes?
  - Use Inter-Process Communication (IPC).

# Linux Signals

- One-way asynchronous notifications sent from:
  - Kernel to process.
  - One process to another process (IPC).
  - A process to itself
- Processes setup signal handlers to control how to respond to a signal
  - Example Ctrl->C or SIGINT to stop a process.

# Linux Signals

- Signal handlers must use signal safe functions which are safe to call asynchronously.
  - Use of global variables can introduce unsafe scenarios.
  - Process code can be interrupted at any time by signals.

# Linux Interprocess Communication (IPC)



Allows process to exchange information without using a common global memory space.

- Pipes
- Semaphores
- Message Queues
- Shared Memory

# Users and Groups

**Advanced Embedded Linux  
Development  
with Dan Walkes**



University of Colorado **Boulder**

**Learning objectives:**  
**Understand Linux Users and Groups**  
**Understand Linux Permissions**



# Linux Users and Groups

- User logs in with a username and password
  - Authentication
- Each user is associated with a User ID (uid)
  - Each process is associated with the UID of the user running the process.
  - `/etc/passwd` maps usernames uids



# Linux Users and Groups

- UID 0 is the root user
  - Can do almost anything on the system
- Each user belongs to one or more groups, with corresponding Group IDs (gid)
  - /etc/group maps group names to gids

# Linux Permissions

Associates a file with owning user, owning group, and permission bits

	User (owner)			Group			Others (everyone)		
	Read (r)	Write (w)	Execute (x)	Read (r)	Write (w)	Execute (x)	Read (r)	Write (w)	Execute (x)
Bit	8	7	6	5	4	3	2	1	0
Octal Value	0400	0200	0100	040	020	010	04	02	01

```
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures$ ls -l
total 24
-rw-r--r-- 1 dan dan 1083 Jun 21 15:27 LICENSE
-rw-r--r-- 1 dan dan 149 Jun 21 15:27 README.md
drwxr-xr-x 4 dan dan 4096 Jun 21 17:08 lecture2
drwxr-xr-x 2 dan dan 4096 Jun 21 15:27 lecture5
drwxr-xr-x 2 dan dan 4096 Jun 21 15:27 lecture7
drwxr-xr-x 4 dan dan 4096 Jun 21 15:27 lecture9
```

# System Programming and Error Handling

**Advanced Embedded Linux  
Development  
with Dan Walkes**



University of Colorado **Boulder**

**Learning objectives:**

**Understand errno.**

**Understand error handling strategies  
for System Programming.**



# Errno and error handling

- C Library mechanism for reporting errors is errno  
`<errno.h>`

# Errno and error handling

- See <https://pubs.opengroup.org/onlinepubs/009695399/basedefs/errno.h.html> for a list of POSIX defined errors - ENOENT means “No such file or directory”
- Use `errno -l` from the command line to dump error values and names

```
dan@DESKTOP-BQMVP69:~/CU/aesd-lectures/lecture2$ errno -l
EPERM 1 Operation not permitted
ENOENT 2 No such file or directory
ESRCH 3 No such process
EINTR 4 Interrupted system call
EIO 5 Input/output error
ENXIO 6 No such device or address
E2BIG 7 Argument list too long
ENOEXEC 8 Exec format error
```

# Errno and error handling

```
int main () {
    const char *filename = "non-existing-file.txt";
    FILE *file = fopen (filename, "rb");
    if (file == NULL) {
        fprintf(stderr, "Value of errno attempting to open file %s: %d\n", filename, errno);
        perror("perror returned");
        fprintf(stderr, "Error opening file %s: %s\n",filename, strerror( errno ));
    } else {
        fclose(file);
    }
    return 0;
}
```

```
Value of errno attempting to open file non-existing-file.txt: 2
perror returned: No such file or directory
Error opening file non-existing-file.txt: No such file or directory
```

# Errno and error handling

PERROR(3)

Linux Programmer's Manual

PERROR(3)

**NAME**`perror - print a system error message`**DESCRIPTION**

The `perror()` function produces a message on standard error describing the last error encountered during a call to a system or library function.

```
int main () {
    const char *filename = "first-non-existing-file.txt";
    FILE *file1 = fopen (filename, "rb");
    const char *filename = "non-existing-file.txt";
    FILE *file = fopen (filename, "rb");
    if (file == NULL) {
        if( file1 ) {
            fclose(file1);
        }
        fprintf(stderr, "Value of errno attempting to open file %s: %d\n", filename, errno);
        perror("perror returned");
        fprintf(stderr, "Error opening file %s: %s\n",filename, strerror( errno ));
    } else {
        fclose(file);
    }
    return 0;
}
```

- What's wrong with this code?
  - `fprintf` may modify `errno`

# Errno and error handling

- How could you fix it?
  - Move fprintf after perror
  - Save errno to local var

```
int main () {
    const char *filename = "first-non-existing-file.txt";
    FILE *file1 = fopen (filename, "rb");
    const char *filename = "non-existing-file.txt";
    FILE *file = fopen (filename, "rb");
    if (file == NULL) {
        if( file1 ) {
            fclose(file1);
        }
        fprintf(stderr, "Value of errno attempting to open file %s: %d\n", filename, errno);
        perror("perror returned");
        fprintf(stderr, "Error opening file %s: %s\n",filename, strerror( errno ));
    } else {
        fclose(file);
    }
    return 0;
}
```



# Errno and error handling

- Given threads use the same memory space
- Wouldn't an error in one thread override an error in a second thread?
  - This is handled for us by POSIX

## Redefinition of `errno`

In POSIX.1, `errno` is defined as an external global variable. But this definition is unacceptable in a multithreaded environment, because its use can result in nondeterministic results. The problem is that two or more threads can encounter errors, all causing the same `errno` to be set. Under these circumstances, a thread might end up checking `errno` after it has already been updated by another thread.

To circumvent the resulting nondeterminism, POSIX.1c redefines `errno` as a service that can access the per-thread error number as follows (ISO/IEC 9945:1-1996, §2.4):

Some functions may provide the error number in a variable accessed through the symbol `errno`. The symbol `errno` is defined by including the header `<errno.h>`, as specified by the C Standard ... For each thread of a process, the value of `errno` shall not be affected by function calls or assignments to `errno` by other threads.

# Embedded Linux Toolchain Overview

**Advanced Embedded Linux  
Development  
with Dan Walkes**



University of Colorado **Boulder**

## **Learning objectives:**

**Understand the components of an  
Embedded Linux toolchain.**

**Understand how to setup an  
Embedded Linux toolchain.**

**Understand the toolchain sysroot.**

# Toolchains

- Compiler
- Linker
- Run-time libraries
- GCC or Clang are the most likely toolchain options.

# GCC Toolchain Components

- Binutils - binary utilities including assembler and linker.
- GCC (Gnu Compiler Collection)
  - Compilers for programming languages (C in our case.)
- C Library - API based on POSIX definition.



# Setting Up a Toolchain

- Option 1: Do it manually by downloading/building/installing components yourself.
- Option 2: Use a build system to generate (for instance Buildroot or Yocto.)

# Setting Up a Toolchain

- Why hand download/build your own Toolchain?
  - You probably shouldn't.
  - We are going to use this with Assignment 2 just to understand how it can be done.
    - Demystify the process, help with build system troubleshooting later.

# Types of Toolchains

- Native toolchain
  - Runs on the same system as the program it generates.
- Cross toolchain
  - Runs on a different architecture than the host.
  - “cross compiling”
    - Creating output for a different hardware architecture.



# Types of Toolchains

- Why use a cross toolchain?
  - Your host is probably more powerful than the target, builds are faster.
  - You probably don't want to include development tools in your target image.
    - Might not be possible to fit these in your target image.

# Specifying Toolchain Targets

- For GNU, a prefix specifies the toolchain target
- Example for QEMU: aarch64-none-linux-gnu-gcc
  - CPU is ARM 64 bit
  - Vendor is “none” (support a common set of ARM CPUs)
  - Kernel is “linux”
  - Operating system is GNU GCC



# Building a toolchain with ARM cross compiler

- See toolchain install instructions at  
<https://github.com/cu-ecen-aeld/aesd-assignments/wiki/Installing-an-ARM-aarch64-developer-toolchain>



# Example Cross Compile Steps

```
yocto@yocto-Latitude-E6540:~/aesd/assignment-3-dwalkes-1/finder-app$ export PATH=$PATH:/usr/local/arm-cross-compiler/install/gcc-arm-10.2-2020.11-x86_64-aarch64-none-linux-gnu/bin
yocto@yocto-Latitude-E6540:~/aesd/assignment-3-dwalkes-1/finder-app$ aarch64-none-linux-gnu-gcc -g -Wall -c -o writer.o writer.c
yocto@yocto-Latitude-E6540:~/aesd/assignment-3-dwalkes-1/finder-app$ file writer.o
writer.o: ELF 64-bit LSB relocatable, ARM aarch64, version 1 (SYSV), with debug_info, not stripped
```

# Sysroot, library and header files

```
yocto@yocto-Latitude-E6540:~/aesd/assignment-3-dwalkes-1/finder-app$ aarch64-none-linux-gnu-gcc -print-sysroot  
/usr/local/arm-cross-compiler/install/gcc-arm-10.2-2020.11-x86_64-aarch64-none-linux-gnu/bin/../aarch64-none-linux-gnu/libc
```

- Sysroot is the root filesystem of your (cross) toolchain.
- Consists of files specific to the \*target\* type.
  - Mirrors files on your host root filesystem.
- Some files are needed to compile programs.
- Others are (also) needed on the target at runtime.



# Sysroot Directories

- lib - Shared objects for C library (on target.)
- usr/lib - Static library archive files for the C library.
- usr/include - Headers for libraries (for instance <stdio.h>.)
- usr/(s)bin: Utility programs for the cross toolchain.

# Sysroot library files

- What do we mean by runtime target files?

```
yocto@yocto-Latitude-E6540:~/aesd/assignment-3-dwalkes-1/finder-app$ file /usr/local/arm-cross-compiler/install/gcc-arm-10.2-2020.11-x86_64-aarch64-none-linux-gnu/bin/..../aarch64-none-linux-gnu/libc/lib64/ld-2.31.so  
/usr/local/arm-cross-compiler/install/gcc-arm-10.2-2020.11-x86_64-aarch64-none-linux-gnu/bin/..../aarch64-none-linux-gnu/libc/lib64/ld-2.31.so: ELF 64-bit LSB shared object, ARM aarch64, version 1 (SYSV), dynamically linked, with debug_info, not stripped
```

- sysroot/lib contains library files specific to the toolchain which will be installed on the target (in this case aarch64) even when built on a different architecture.

# Toolchain Sysroot files

- Will the aarch64-none-linux-gnu-gcc compiler be in the toolchain sysroot?
  - No, because we are cross compiling
  - The aarch64-none-linux-gnu-gcc compiler is in your host system filesystem (and run by your host system), not your toolchain system root

# Other tools in the toolchain

- Use all cross toolchain components with the same prefix referenced for gcc (aarch64-none-linux-gnu-XXXX)
- gcc, g++ - compiler
- gdb - debugger
- ld - linker



# Other tools in the toolchain

- `addr2line` : Converts program addresses into filenames/numbers for debug.
- `objdump` - Disassemble object files.
- `strip` - Remove debug tables, make binary files smaller.
- `readelf` - Additional information about executables (object code, location in memory map, etc).

# Static vs Dynamic Linking

- gcc or g++ always links with glibc, the C library
- Static Linkage
  - All library functions and dependencies are pulled from archive and placed in your executable
- Dynamic Linkage
  - Linking is done dynamically at runtime

# Static vs Dynamic Linking

- When to use Static Linkage?
  - When you have relatively few applications (or only a single application)
    - Busybox
  - You need to run an application before the root filesystem is available
    - When would that happen?
      - At boot, loading storage drivers.

# Shared Libraries (Dynamic Link)

```
yocto@yocto-Latitude-E6540:~/aesd/assignment-3-dwalkes-1/finder-app$ make CROSS_COMPILE=aarch64-none-linux-gnu- clean  
rm -f *.o writer *.elf *.map  
yocto@yocto-Latitude-E6540:~/aesd/assignment-3-dwalkes-1/finder-app$ make CROSS_COMPILE=aarch64-none-linux-gnu- all  
aarch64-none-linux-gnu-gcc -g -Wall -c -o writer.o writer.c  
aarch64-none-linux-gnu-gcc -g -Wall -I/ writer.o -o writer  
yocto@yocto-Latitude-E6540:~/aesd/assignment-3-dwalkes-1/finder-app$ aarch64-none-linux-gnu-readelf -a writer | grep "Shared library"  
 0x0000000000000001 (NEEDED)           Shared library: [libc.so.6]  
yocto@yocto-Latitude-E6540:~/aesd/assignment-3-dwalkes-1/finder-app$  
  
yocto@yocto-Latitude-E6540:~/aesd/assignment-3-dwalkes-1/finder-app$ aarch64-none-linux-gnu-readelf -a writer | grep "program interpreter"  
[Requesting program interpreter: /lib/ld-linux-aarch64.so.1]
```

- Why is libc/ld-linux-aarch64 listed here?
  - This is a shared libraries referenced by “writer”
- What does it mean that libc is listed here?
  - This file need to be available on the root filesystem to run “writer” successfully.



# Shared Library Locations

- Linker checks for shared libraries in:
  - /lib, /lib64
  - /usr/lib, /usr/lib64
  - content of LD\_LIBRARY\_PATH

# Logging and Syslog

**Advanced Embedded Linux  
Development  
with Dan Walkes**



University of Colorado **Boulder**

**Learning objectives:**  
**Logging and Syslog Overview.**



# Printf Debugging

- Print output is one of the simplest ways to troubleshoot and debug your application.
- The printf call works great for interactive programs.
- What about headless embedded systems?



# Logging

- What about headless embedded systems?
  - Need some type of logging framework to store information in log files.
  - Syslog and syslogd (rsyslogd) is one such framework.



# Logging and Syslog

- `syslogd` is a daemon which
  - Uses configuration file to configure logging (usually writing to files in `/var/log`)
  - Handles log messages from applications using `syslog()` API calls.

# Syslog Usage

## SYNOPSIS

```
#include <syslog.h>

void openlog(const char *ident, int option, int facility);
void syslog(int priority, const char *format, ...);
```

- Facilities
  - LOG\_USER
  - LOG\_DAEMON
  - LOG\_LOCAL0 to LOG\_LOCAL7
- Priorities
  - LOG\_ERR
  - LOG\_WARNING
  - LOG\_INFO
  - LOG\_DEBUG

# Syslog Usage

```
openlog(NULL,0,LOG_USER);
```

```
syslog(LOG_ERR,"Invalid Number of arguments: %d",argc);
```

- Results in message in /var/log/syslog:

```
|Sep  2 08:36:01 ecen5013-VirtualBox anacron[825]: Normal exit (3 jobs run)
|Sep  2 08:45:41 ecen5013-VirtualBox writer: Invalid Number of arguments: 1
|...
```

# Syslog Usage

- Log redirection in an Ubuntu VM happens based on rules in  
`/etc/rsyslog.d/*default.conf`

```
# Default rules for rsyslog.
#
# For more information see rsyslog.conf(5) and /etc/rsyslog.conf
#
# First some standard log files. Log by facility.
#
auth,authpriv.*          /var/log/auth.log
*.*;auth,authpriv.none   -/var/log/syslog
#cron.*                  /var/log/cron.log
#daemon.*                -/var/log/daemon.log
kern.*                   -/var/log/kern.log
#lpr.*                   -/var/log/lpr.log
mail.*                   -/var/log/mail.log
#user.*                  -/var/log/user.log
"
```